

Go, The Standard Library

Real Code. Real Productivity.
Master The Go Standard Library

Daniel Huckstep



Go, The Standard Library

Real Code. Real Productivity. Master The Go Standard Library

Daniel Huckstep

This book is for sale at <http://leanpub.com/go-thestdlib>

This version was published on 2017-03-26



This is a [Leanpub](#) book. Leanpub empowers authors and publishers with the Lean Publishing process. [Lean Publishing](#) is the act of publishing an in-progress ebook using lightweight tools and many iterations to get reader feedback, pivot until you have the right book and build traction once you do.

© 2012 - 2017 Daniel Huckstep

Tweet This Book!

Please help Daniel Huckstep by spreading the word about this book on [Twitter](#)!

The suggested hashtag for this book is [#GoTheStdLib](#).

Find out what other people are saying about the book by clicking on this link to search for this hashtag on Twitter:

<https://twitter.com/search?q=#GoTheStdLib>

Contents

Introduction	1
Target Audience	1
How To Read This Book	2
Code In The Book	2
Thanks	5
archive	6
Meet The Archive Package	6
Writing tar Files	6
Writing zip Files	9
Reading tar Files	12
Reading zip Files	15
Caveats	16
builtin	18
Batteries Included	18
Building Objects	18
Maps, Slices, And Channels	20
All The Sizes	24
Causing And Handling Panics	26
Complex Numbers	28
expvar	29

Introduction

When I sit down to build a new piece of software in my favorite programming language of the week, I open up my programmer's toolbox. I can pull out a number of things, like my knowledge of the language syntax and its quirks. It probably has some sort of library packaging system ([rubygems](http://rubygems.org/)¹ or [python eggs](http://pypi.python.org/pypi/)²), and I have my list of libraries for doing certain jobs. The language also has a **standard library**. All of these tools combine to help solve difficult programming problems.

Right now, my programming language of choice is [Go](http://golang.org/)³ and it has a wonderful standard library. That standard library is what this book is about.

I wanted to take an in depth look at something which normally doesn't get a lot of press, and many developers overlook. The standard library usually has a number of great solutions to problems that you might be using some other dependency for, simply because you don't know about them. *It makes no sense for my application to depend on an external library or program if the standard distribution of the language has something built in.*⁴

Learning the ins and outs of your favorite programming language's standard library can help make you a better programmer, and streamline your applications by removing dependencies. If this sounds like something you're interested in, keep reading.

Target Audience

This book is for people that know how to program Go already. It's definitely not an intro. If you're completely new to Go, start with [the documentation page](#)⁵ and [the reference page](#)⁶. The language specification is quite readable and if you're already familiar with other programming languages you can probably absorb the language from the spec.

If you know Go but want to step up your game and your usage of the standard library, this book is for you.

¹<http://rubygems.org/>

²<http://pypi.python.org/pypi/>

³<http://golang.org/>

⁴Not to mention, the library you are using might only work on one operating system, while the standard library should work everywhere the language works.

⁵<http://golang.org/doc/>

⁶<http://golang.org/ref/>

How To Read This Book

My goal for this book is a *readable reference*. I do want you to read it, but I also want you to be able to pull it off the electronic shelf and remind yourself of how to do something, like writing a zip file. It's not meant to be a replacement for [the package reference](http://golang.org/pkg/)⁷ which is very useful to remember the details about a specific method/function/type/interface.

So feel free to read from cover to cover, and in fact I recommend this approach. If you see something that doesn't quite work reading it this way, let me know. Alternatively, try reading individual chapters when you start to deal with a given package to get a feel for it, and come back to skim to refresh your memory.

Code In The Book

All the code listed in the book is available for download from Leanpub as an extra. Visit your [dashboard](https://leanpub.com/dashboard)⁸ for access to the archives.

Anything with a main package should be able to be executed with `go run` by Go Version 1.2. If it's not, please let me know, with as much error information as possible.

Some code may depend on output from previously shown code in the same chapter. For example, the tar archive reading code reads the tar created in the writing code.

Frequently I'll use other packages to make my life easier when writing example code. Don't worry too much about it. If you're confused about some use of a package you're not familiar with yet, either try to ignore the details and trust that I'll explain it later, or jump ahead and choose your own adventure!

License

Code distributed as part of this book, either inline or with the above linked archive, is licensed under the MIT license:

⁷<http://golang.org/pkg/>

⁸<https://leanpub.com/dashboard>

LICENSE

```
1 Copyright (c) 2014 Daniel Huckstep
2
3 Permission is hereby granted, free of charge, to any person obtaining a copy of \
4 this software and associated documentation files (the "Software"), to deal in th\
5 e Software without restriction, including without limitation the rights to use, \
6 copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the \
7 Software, and to permit persons to whom the Software is furnished to do so, subj\
8 ect to the following conditions:
9
10 The above copyright notice and this permission notice shall be included in all c\
11 opies or substantial portions of the Software.
12
13 THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLI\
14 ED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR \
15 A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYR\
16 IGH T HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN \
17 ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WIT\
18 H THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.
```

Some code is taken directly from the Go source distribution. This code is licensed under a BSD-style license by The Go Authors:

GOLICENSE

```
1 Copyright (c) 2012 The Go Authors. All rights reserved.
2
3 Redistribution and use in source and binary forms, with or without
4 modification, are permitted provided that the following conditions are
5 met:
6
7     * Redistributions of source code must retain the above copyright
8 notice, this list of conditions and the following disclaimer.
9     * Redistributions in binary form must reproduce the above
10 copyright notice, this list of conditions and the following disclaimer
11 in the documentation and/or other materials provided with the
12 distribution.
13     * Neither the name of Google Inc. nor the names of its
14 contributors may be used to endorse or promote products derived from
15 this software without specific prior written permission.
16
17 THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS
```

18 "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT
19 LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR
20 A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
21 OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
22 SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
23 LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
24 DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
25 THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
26 (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
27 OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Thanks

Thanks for buying and checking out this book. As part of the lean publishing philosophy, you'll be able to interact with me as the book is completed. I'll be able to change things, reorganize parts, and generally make a better book. I hope you enjoy.

A big thanks goes out to all those who provided feedback during the writing process:

- Brad Fitzpatrick
- Mikhail Strebkov
- Kim Shrier

archive

Meet The Archive Package

The `archive` package is used to read and write files in tar and zip format. Both formats pack multiple files into one big file, the main difference being that zip files support optional compression using the DEFLATE algorithm provided by the `compress/flate` package.

Writing tar Files

Writing a tar file starts with `NewWriter`. It takes an `io.Writer` type, which is just something that has a method that looks like `Write([]byte) (int, error)`. This is nice if you want to generate a tar file on the fly and write it out to an HTTP response, or feed it through another writer like a gzip writer. You'll see this *just give me an `io.Writer`* pattern a lot in the Go stdlib. In our case, I'm just going to write the archive out to a file.



Make sure to close the writer you pass in *after* you close the tar writer.

It writes 2 zero blocks to finish up the file, but ignores any errors during this process. This *trailer* isn't strictly required, but it's good to have. If you use `defer` in the natural order, you should be okay.

To add files to the new tar writer, use `WriteHeader`. It needs a `Header` with all the information about this entry in the archive, including its name, size, permissions, user and group information, and all the other bits that get set when the tar file gets unpacked. Straight from the Go documentation, the `Header` type looks like this:

archive/tar_header.go

```
1 type Header struct {
2     Name      string    // name of header file entry
3     Mode      int64     // permission and mode bits
4     Uid       int      // user id of owner
5     Gid       int      // group id of owner
6     Size      int64     // length in bytes
7     ModTime   time.Time // modified time
8     Typeflag   byte     // type of header entry
9     Linkname   string    // target name of link
10    Uname      string    // user name of owner
11    Gname      string    // group name of owner
12    Devmajor   int64     // major number of character or block device
13    Devminor   int64     // minor number of character or block device
14    AccessTime time.Time // access time
15    ChangeTime time.Time // status change time
16 }
```

Some fields aren't really required if you're doing something quick and dirty, and some only apply to certain types of entries (controlled by the `Typeflag` field). For example, if you're packaging a regular file, you don't need to worry about `Devmajor` and `Devminor`.



I found that on top of the obvious `Name` and `Size` fields, I had to set the `ModTime` on the `Header`. GNU tar would unpack the file fine, but running the read script would throw the standard “archive/tar: invalid tar header” error back at me.

Let's see it all together:

archive/write_tar.go

```
1 package main
2
3 import (
4     "archive/tar"
5     "fmt"
6     "io"
7     "log"
8     "os"
9 )
10
```

```
11 var files = []string{"write_tar.go", "read_tar.go"}
12
13 func addFile(filename string, tw *tar.Writer) error {
14     file, err := os.Open(filename)
15     if err != nil {
16         return fmt.Errorf("failed opening %s: %s", filename, err)
17     }
18     defer file.Close()
19
20     stat, err := file.Stat()
21     if err != nil {
22         return fmt.Errorf("failed file stat for %s: %s", filename, err)
23     }
24
25     hdr := &tar.Header{
26         ModTime: stat.ModTime(),
27         Name:     filename,
28         Size:     stat.Size(),
29         Mode:     int64(stat.Mode().Perm()),
30     }
31
32     if err := tw.WriteHeader(hdr); err != nil {
33         msg := "failed writing tar header for %s: %s"
34         return fmt.Errorf(msg, filename, err)
35     }
36
37     copied, err := io.Copy(tw, file)
38     if err != nil {
39         return fmt.Errorf("failed writing %s to tar: %s", filename, err)
40     }
41
42     // Check copied, since we have the file stat with its size
43     if copied < stat.Size() {
44         msg := "wrote %d bytes of %s, expected to write %d"
45         return fmt.Errorf(msg, copied, filename, stat.Size())
46     }
47
48     return nil
49 }
50
51 func main() {
52     flags := os.O_WRONLY | os.O_CREATE | os.O_TRUNC
```

```
53     file, err := os.OpenFile("go.tar", flags, 0644)
54     if err != nil {
55         log.Fatalf("failed opening tar for writing: %s", err)
56     }
57     defer file.Close()
58
59     tw := tar.NewWriter(file)
60     defer tw.Close()
61
62     for _, filename := range files {
63         if err := addFile(filename, tw); err != nil {
64             log.Fatalf("failed adding file %s to tar: %s", filename, err)
65         }
66     }
67 }
```

Remember to `Close` the tar writer first, followed by the original `io.Writer`. In the example, I defer the calls to `Close`. Because `defer` executes in a LIFO^a order, this is exactly the order things get closed in. `defer` usually results in you not having to think too hard in these situations, just use `defer` the way it should be used, and everything should be fine.

^aLast In First Out

Writing zip Files

Writing a zip file is similar to writing a tar file. There's a `NewWriter` function that takes an `io.Writer`, so let's use that.

The zip package has a handy helper to let you quickly write a file to the archive without much ceremony. We can use the `Create(name string)` method on the zip writer we got back from `NewWriter` to add an entry to the zip; no header information needed. There is a `Header` type, which looks like this:

archive/zip_header.go

```
1 type FileHeader struct {
2     Name          string
3     CreatorVersion uint16
4     ReaderVersion  uint16
5     Flags          uint16
6     Method         uint16
7     ModifiedTime   uint16 // MS-DOS time
8     ModifiedDate   uint16 // MS-DOS date
9     CRC32          uint32
10    CompressedSize  uint32 // deprecated; use CompressedSize64
11    UncompressedSize uint32 // deprecated; use UncompressedSize64
12    CompressedSize64 uint64
13    UncompressedSize64 uint64
14    Extra           []byte
15    ExternalAttrs    uint32 // Meaning depends on CreatorVersion
16    Comment          string
17 }
```

You *can* use `CreateHeader` if you need to do something special, but `Create` creates a basic header for us and gives us a writer back. We can now use this writer to write the file into the zip archive.

Make sure to write the entire file before calling any of `Create`, `CreateHeader`, or `Close`. You can only deal with one file at a time, and you certainly can't deal with the zip after you've closed it.

archive/write_zip.go

```
1 package main
2
3 import (
4     "archive/zip"
5     "fmt"
6     "io"
7     "log"
8     "os"
9 )
10
11 var files = []string{"write_zip.go", "read_zip.go"}
12
13 func addFile(filename string, zw *zip.Writer) error {
```

```

14     file, err := os.Open(filename)
15     if err != nil {
16         return fmt.Errorf("failed opening %s: %s", filename, err)
17     }
18     defer file.Close()
19
20     wr, err := zw.Create(filename)
21     if err != nil {
22         msg := "failed creating entry for %s in zip file: %s"
23         return fmt.Errorf(msg, filename, err)
24     }
25
26     // Not checking how many bytes copied,
27     // since we don't know the file size without doing more work
28     if _, err := io.Copy(wr, file); err != nil {
29         return fmt.Errorf("failed writing %s to zip: %s", filename, err)
30     }
31
32     return nil
33 }
34
35 func main() {
36     flags := os.O_WRONLY | os.O_CREATE | os.O_TRUNC
37     file, err := os.OpenFile("go.zip", flags, 0644)
38     if err != nil {
39         log.Fatalf("failed opening zip for writing: %s", err)
40     }
41     defer file.Close()
42
43     zw := zip.NewWriter(file)
44     defer zw.Close()
45
46     for _, filename := range files {
47         if err := addFile(filename, zw); err != nil {
48             log.Fatalf("failed adding file %s to zip: %s", filename, err)
49         }
50     }
51 }

```

As with tar files, remember to `Close` the original `io.Writer` and the zip writer (in that order).

Reading tar Files

Reading tar files is pretty straight forward. You use `NewReader` to get a handle to a `Reader` type. Like `NewWriter` taking an `io.Writer` type, `NewReader` takes an `io.Reader` type, in order to plug into other streams for reading tar files on the fly.

Once you have your `Reader`, you can iterate over the entries in the archive with the `Next` method. It returns a `Header` and possibly an `error`. Remember to check the error since it's used to signal the end of the archive (with `io.EOF`) and other problems.

Always check those errors!

You can read out an entry by calling `Read` on the reader you got back from `NewReader`, or pass it to a utility function to read out the full contents of the entry. In the example, I use `io.ReadFull` to read out the appropriate number of bytes into a slice, and can then print that to `stdout`.

archive/read_tar.go

```
1 package main
2
3 import (
4     "archive/tar"
5     "fmt"
6     "io"
7     "log"
8     "os"
9     "text/template"
10 )
11
12 var HeaderTemplate = `tar header
13 Name:      {{.Name}}
14 Mode:      {{.Mode | printf "%o" }}
15 UID:       {{.Uid}}
16 GID:       {{.Gid}}
17 Size:      {{.Size}}
18 ModTime:   {{.ModTime}}
19 Typeflag:  {{.Typeflag | printf "%q" }}
20 Linkname:  {{.Linkname}}
21 Uname:     {{.Uname}}
22 Gname:     {{.Gname}}
23 Devmajor:  {{.Devmajor}}
24 Devminor:  {{.Devminor}}
25 AccessTime: {{.AccessTime}}
26 ChangeTime: {{.ChangeTime}}
```

```
27 `
28 var CompiledHeaderTemplate *template.Template
29
30 func init() {
31     t := template.New("header")
32     CompiledHeaderTemplate = template.Must(t.Parse(HeaderTemplate))
33 }
34
35 func printHeader(hdr *tar.Header) {
36     CompiledHeaderTemplate.Execute(os.Stdout, hdr)
37 }
38
39 func printContents(tr io.Reader, size int64) {
40     contents := make([]byte, size)
41     read, err := io.ReadFull(tr, contents)
42
43     if err != nil {
44         log.Fatalf("failed reading tar entry: %s", err)
45     }
46
47     if int64(read) != size {
48         log.Fatalf("read %d bytes but expected to read %d", read, size)
49     }
50
51     fmt.Fprintf(os.Stdout, "Contents:\n\n%s", contents)
52 }
53
54 func main() {
55     file, err := os.Open("go.tar")
56     if err != nil {
57         msg := "failed opening archive, run `go run write_tar.go` first: %s"
58         log.Fatalf(msg, err)
59     }
60
61     defer file.Close()
62
63     tr := tar.NewReader(file)
64     for {
65         hdr, err := tr.Next()
66         if err == io.EOF {
67             break
68         }
69     }
```

```
69
70         if err != nil {
71             log.Fatalf("failed getting next tar entry: %s", err)
72         }
73
74         printHeader(hdr)
75         printContents(tr, hdr.Size)
76     }
77 }
```

Output:

```
1 tar header
2 Name:      write_tar.go
3 Mode:      644
4 UID:       0
5 GID:       0
6 Size:      1441
7 ModTime:   2014-03-07 23:02:17 -0700 MST
8 Typeflag:  '\x00'
9 Linkname:
10 Uname:
11 Gname:
12 Devmajor:  0
13 Devminor:  0
14 AccessTime: 0001-01-01 00:00:00 +0000 UTC
15 ChangeTime: 0001-01-01 00:00:00 +0000 UTC
16 Contents:
17
18 <snip contents of writer_tar.go>
19 tar header
20 Name:      read_tar.go
21 Mode:      644
22 UID:       0
23 GID:       0
24 Size:      1484
25 ModTime:   2014-03-07 23:00:03 -0700 MST
26 Typeflag:  '\x00'
27 Linkname:
28 Uname:
29 Gname:
30 Devmajor:  0
```

```
31 Devminor: 0
32 AccessTime: 0001-01-01 00:00:00 +0000 UTC
33 ChangeTime: 0001-01-01 00:00:00 +0000 UTC
34 Contents:
35
36 <snip contents of read_tar.go>
```

Reading zip Files

Reading zip files is a walk in the park too. Start with `OpenReader` to get a `zip.ReadCloser`. It has a collection of `File` structs you can iterate through, each one with size and other information, and an `Open` method so you can get another `ReadCloser` to read out that individual file. Simple!

archive/read_zip.go

```
1 package main
2
3 import (
4     "archive/zip"
5     "fmt"
6     "io"
7     "log"
8     "os"
9 )
10
11 func printFile(file *zip.File) error {
12     frc, err := file.Open()
13     if err != nil {
14         msg := "failed opening zip entry %s for reading: %s"
15         return fmt.Errorf(msg, file.Name, err)
16     }
17     defer frc.Close()
18
19     fmt.Fprintf(os.Stdout, "Contents of %s:\n", file.Name)
20
21     copied, err := io.Copy(os.Stdout, frc)
22     if err != nil {
23         msg := "failed reading zip entry %s for reading: %s"
24         return fmt.Errorf(msg, file.Name, err)
25     }
26 }
```

```
27         if uint64(copied) != file.UncompressedSize64 {
28             msg := "read %d bytes of %s but expected to read %d bytes"
29             return fmt.Errorf(msg, copied, file.UncompressedSize64)
30         }
31
32         fmt.Println()
33
34         return nil
35     }
36
37     func main() {
38         rc, err := zip.OpenReader("go.zip")
39         if err != nil {
40             msg := "failed opening archive, run `go run write_zip.go` first: %s"
41             log.Fatalf(msg, err)
42         }
43         defer rc.Close()
44
45         for _, file := range rc.File {
46             if err := printFile(file); err != nil {
47                 log.Fatalf("failed reading %s from zip: %s", file.Name, err)
48             }
49         }
50     }
```

Output:

```
1 Contents of write_zip.go:
2 <snip contents of write_zip.go>
3
4 Contents of read_zip.go:
5 <snip contents of read_zip.go>
```

Remember to `Close` the first `ReadCloser` you get from `OpenReader`, as well as all the other ones you get while reading files.

Caveats

ZIP64

You may have noticed the `FileHeader` has two pairs of numbers for the size of a file in the archive. The `CompressedSize` and `UncompressedSize` are `uint32` values. These

are deprecated, but in the interest of backwards compatibility will still work for regular zip files. If you're working with ZIP64 files, you need to use the newer `CompressedSize64` and `UncompressedSize64` `uint64` values. These will be correct for all files, so they are the preferred values to use.

builtin

Batteries Included

The `builtin` package isn't a real package, it's just here to document the builtin functions that come with the language. Lower level than the standard library, these things are just...there. The builtins let you do things with maps, slices, channels, and imaginary numbers, cause and deal with panics, build objects, and get size information about certain things. Honestly, most of this can be learned from the spec, but I've included it for completeness.

Building Objects

make

`make` is used to build the builtin types like slices, channels and maps. The first argument is the type, and it can be one of those three types.

In the case of channels, there is an optional second integer parameter, the *capacity*. If it's zero (or not given), the channel is unbuffered. This means writes block until there is a reader ready to receive the data, and reads block until there is a write ready to give data. If the parameter is greater than zero, the channel is buffered with the capacity specified. On these channels, reads block only when the channel is empty, and writes block only when the channel is full.

In the case of maps, the second parameter is also optional, but is rarely used. It controls the initial allocation, so if you know exactly how big your map has to be, it can be helpful. `cap` (which we'll see later) doesn't work on maps though, so you can't really examine the effects of this second parameter easily.

In the case of slices, the second parameter is **not** optional, and specifies the starting length of the slice. Oh but the plot thickens! There is an optional third parameter, which controls the starting capacity, and it can't be smaller than the length.⁹ This way, you can get really specific with your slice allocation and save subsequent reallocations if you know exactly how much space you need it to take up.

⁹If you specify a length greater than the capacity, you'll get a runtime panic.

builtin/make.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     unbuffered := make(chan int)
7     log.Printf("unbuffered: %v, type: %T, len: %d, cap: %d", unbuffered, unbuffered\
8 , len(unbuffered), cap(unbuffered))
9
10    buffered := make(chan int, 10)
11    log.Printf("buffered: %v, type: %T, len: %d, cap: %d", buffered, buffered, len(\
12 buffered), cap(buffered))
13
14    m := make(map[string]int)
15    log.Printf("m: %v, len: %d", m, len(m))
16
17    // Would cause a compile error
18    // slice := make([]byte)
19
20    slice := make([]byte, 5)
21    log.Printf("slice: %v, len: %d, cap: %d", slice, len(slice), cap(slice))
22
23    slice2 := make([]byte, 0, 10)
24    log.Printf("slice: %v, len: %d, cap: %d", slice2, len(slice2), cap(slice2))
25 }
```

new

The `new` function allocates a new object of the type provided, and returns a pointer to the new object. The object is allocated to be the zero value for the given type. It's not something you use terribly often, but it can be useful. If you're making a new struct, you probably want to use the composite literal syntax instead.

builtin/new.go

```
1 package main
2
3 import "log"
4
5 type Actor struct {
6     Name string
7 }
8
9 type Movie struct {
10     Title string
11     Actors []*Actor
12 }
13
14 func main() {
15     ip := new(int)
16     log.Printf("ip type: %T, ip: %v, *ip: %v", ip, ip, *ip)
17
18     m := new(Movie)
19     log.Printf("m type: %T, m: %v, *m: %v", m, m, *m)
20 }
```

Maps, Slices, And Channels

You've got slices, maps and channels as some of the fundamental types that Go provides. The functions `delete`, `close`, `append`, and `copy` all deal with these types to do basic operations.

delete

`delete` removes elements from a map. If the key doesn't exist in the map, nothing happens, nothing to worry about. If the map itself is `nil` it still works, just nothing happens.

builtin/delete.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     m := make(map[string]int)
7     log.Println(m)
8
9     m["one"] = 1
10    log.Println(m)
11
12    m["two"] = 2
13    log.Println(m)
14
15    delete(m, "one")
16    log.Println(m)
17
18    delete(m, "one")
19    log.Println(m)
20
21    m = nil
22    delete(m, "two")
23 }
```

close

`close` takes a writable channel and closes it. When I say writable, I mean either a *normal* channel like `var normal chan int` or a *write only* channel like `var writeOnly chan<- int`. You can still receive from a closed channel, but you'll get the *zero value* of whatever the type is. If you want to check that you actually got a value and not the zero value, use the *comma ok* pattern. Closing an already closed channel will panic, so watch those double closes.

builtin/close.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     c := make(chan int, 1)
7     c <- 1
8
9     log.Println(<-c) // Prints 1
10
11    c <- 2
12    close(c)
13
14    log.Println(<-c) // Prints 2
15    log.Println(<-c) // Prints 0
16
17    if i, ok := <-c; ok {
18        log.Printf("Channel is open, got %d", i)
19    } else {
20        log.Printf("Channel is closed, got %d", i)
21    }
22
23    close(c) // Panics, channel is already closed
24 }
```

append

`append` tacks on elements to the end of a slice, exactly like it sounds. You need to keep the return value around, since it's the new slice with the extra data. It could return the same slice if it has space for the data, but it might return something new if it needed to allocate more memory. It takes a variable number of arguments, so if you want to append an existing array, use `...` to expand the array.

The idiomatic way to append to a slice is to assign the result to the same slice you're appending to. It's probably what you want.

builtin/append.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     // Empty slice, with capacity of 10
7     ints := make([]int, 0, 10)
8     log.Printf("ints: %v", ints)
9
10    ints2 := append(ints, 1, 2, 3)
11
12    log.Printf("ints2: %v", ints2)
13    log.Printf("Slice was at %p, it's probably still at %p", ints, ints2)
14
15    moreInts := []int{4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14}
16    ints3 := append(ints2, moreInts...)
17
18    log.Printf("ints3: %v", ints3)
19    log.Printf("Slice was at %p, and it moved to %p", ints2, ints3)
20
21    ints4 := []int{1, 2, 3}
22    log.Printf("ints4: %v", ints4)
23    // The idiomatic way to append to a slice,
24    // just assign to the same variable again
25    ints4 = append(ints4, 4, 5, 6)
26    log.Printf("ints4: %v", ints4)
27 }
```

copy

`copy` copies from one slice to another. It will also copy *from* a string, treating it as a slice of bytes. It returns the number of bytes copied, which is the shorter of the lengths of the two slices.

builtin/copy.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     ints := []int{1, 2, 3, 4, 5, 6}
7     otherInts := []int{11, 12, 13, 14, 15, 16}
8
9     log.Printf("ints: %v", ints)
10    log.Printf("otherInts: %v", otherInts)
11
12    copied := copy(ints[:3], otherInts)
13    log.Printf("Copied %d ints from otherInts to ints", copied)
14
15    log.Printf("ints: %v", ints)
16    log.Printf("otherInts: %v", otherInts)
17
18    hello := "Hello, World!"
19    bytes := make([]byte, len(hello))
20
21    copy(bytes, hello)
22
23    log.Printf("bytes: %v", bytes)
24    log.Printf("hello: %s", hello)
25 }
```

All The Sizes

A lot of things have lengths and capacities. With `len` and `cap`, you can find out about these values.

len

`len` tells you the actual *length* or size of something. In the case of slices, you get, well, the length. In the case of strings, you get the number of bytes. For maps, you get how many pairs are in the map. For channels, you get how many elements the channel has buffered (only relevant for buffered channels).

You can also call `len` with a pointer, but only a pointer to an array. It's the equivalent of calling it on the dereferenced pointer. But, since it still has a type, it's an *array*

and not a *slice*, and the type of an array includes the size, so it still works. The length is part of the type.

builtin/len.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     slice := make([]byte, 10)
7     log.Printf("slice: %d", len(slice))
8
9     str := "γεία σου κόσμο"
10    log.Printf("string: %d", len(str))
11
12    m := make(map[string]int)
13    m["hello"] = 1
14    log.Printf("map: %d", len(m))
15
16    channel := make(chan int, 5)
17    log.Printf("channel: %d", len(channel))
18    channel <- 1
19    log.Printf("channel: %d", len(channel))
20
21    var pointer *[5]byte
22    log.Printf("pointer: %d", len(pointer))
23 }
```

cap

`cap` tells you the capacity of something. It's similar to `len`, except it doesn't work on maps or strings. With arrays, it's the same as using `len`.

With slices, it returns the max size the slice can grow to when you append to it before things are copied to a new backing array. This is why you have to save the return value of `append`. If `cap` returns 5 and you append 6 things to your slice, it's going to return you a slice backed by a new array.

With channels, it returns the buffer capacity.

builtin/cap.go

```

1 package main
2
3 import "log"
4
5 func main() {
6     slice := make([]byte, 0, 5)
7     log.Printf("slice: %d", cap(slice))
8
9     channel := make(chan int, 10)
10    log.Printf("channel: %d", cap(channel))
11
12    var pointer *[15]byte
13    log.Printf("pointer: %d == %d", cap(pointer), len(pointer))
14 }

```

Causing And Handling Panics

`panic` and `recover` are typically used to deal with errors. These are errors where returning an error in the *comma err* style don't make sense. Things like programmer error or things that are seriously broken. *Usually*.

If bad things are afoot, you can use `panic` to throw an error. You can pass it pretty much any object, which gets carried up the stack. Deferred functions get executed, and up the error goes. It works sort of like `raise` or `throw` in other languages.

You can use `recover` to, as the name says, recover from a panic. `recover` must be executed from *within* a deferred function, and not from within a function the deferred function calls. It returns whatever panic was called with, you check for `nil` and can then type cast it to something.



There are some creative uses¹⁰ for `panic/recover` beyond error handling, but they should be confined to your own package. In Go, it's not nice to let a panic go outside your own little world. Better to handle the panic yourself in a way you know how, and return an appropriate error. In some cases, the panic makes sense. Err on the side of returning instead of panicking.

The example illustrates things much better.

¹⁰See the code for the `encoding/json` package on one of them.

builtin/panic_recover.go

```
1 package main
2
3 import (
4     "errors"
5     "log"
6 )
7
8 func handlePanic(f func()) {
9     defer func() {
10         if r := recover(); r != nil {
11             if str, ok := r.(string); ok {
12                 log.Printf("got a string error: %s", str)
13                 return
14             }
15
16             if err, ok := r.(error); ok {
17                 log.Printf("got an error error: %s", err.Error())
18                 return
19             }
20
21             log.Printf("got a different kind of error: %v", r)
22         }
23     }()
24     f()
25 }
26
27 func main() {
28     handlePanic(func() {
29         panic("string error")
30     })
31
32     handlePanic(func() {
33         panic(errors.New("error error"))
34     })
35
36     handlePanic(func() {
37         panic(10)
38     })
39 }
```

Complex Numbers

Go supports complex numbers as a builtin type. You can define them with literal syntax, or by using the builtin function `complex`. If you want to build a complex number from existing float values, you need to use the builtin function, and the two arguments have to be of the same type (`float32` or `float64`) and will produce a complex type double the size (`complex64` or `complex128`). Once you have a complex number, you can add, subtract, divide, and multiply values normally.

If you have a complex number and want to break it into the real and imaginary parts, use the functions `real` and `imag`.

builtin/complex.go

```
1 package main
2
3 import "log"
4
5 func main() {
6     c1 := 1.5 + 0.5i
7     c2 := complex(1.5, 0.5)
8     log.Printf("c1: %v", c1)
9     log.Printf("c2: %v", c2)
10    log.Printf("c1 == c2: %v", c1 == c2)
11    log.Printf("c1 real: %v", real(c1))
12    log.Printf("c1 imag: %v", imag(c1))
13    log.Printf("c1 + c2: %v", c1+c2)
14    log.Printf("c1 - c2: %v", c1-c2)
15    log.Printf("c1 * c2: %v", c1*c2)
16    log.Printf("c1 / c2: %v", c1/c2)
17    log.Printf("c1 type: %T", c1)
18
19    c3 := complex(float32(1.5), float32(0.5))
20    log.Printf("c3 type: %T", c3)
21 }
```

expvar

The `expvar` package is global variables done right.

It has helpers for `Float`, `Int`, `Map`, and `String` types, which are setup to be atomic. Things are registered by a string name, the `Key`, and they map to a corresponding `Var`, which is just an interface with a single method: `String() string`.

This simple interface allows you to use the more raw `Publish` method to register more custom handlers in the form of a `Func` type. These are just functions which take no arguments and return an empty interface (which, in implementation should probably be a string).

Examining the source for the package, you can see it uses this to register the `memstats` variable. When you iterate through the variables and you call the `String` method on the `Var`, the function runs to extract the `memstats` at that moment in time.

It's a pretty simple, but very powerful package. You can use it for metric type stuff, or you can use it as a more traditional global variable system. It can do it all.

expvar/expvar.go

```
1 package main
2
3 import (
4     "expvar"
5     "flag"
6     "log"
7     "time"
8 )
9
10 var (
11     times      = flag.Int("times", 1, "times to say hello")
12     name       = flag.String("name", "World", "thing to say hello to")
13     helloTimes = expvar.NewInt("hello")
14 )
15
16 func init() {
17     expvar.Publish("time", expvar.Func(now))
18 }
19
```

```
20 func now() interface{} {
21     return time.Now().Format(time.RFC3339Nano)
22 }
23
24 func hello(times int, name string) {
25     helloTimes.Add(int64(times))
26     for i := 0; i < times; i++ {
27         log.Printf("Hello, %s!", name)
28     }
29 }
30
31 func printVars() {
32     log.Println("expvars:")
33     expvar.Do(func(kv expvar.KeyValue) {
34         switch kv.Key {
35             case "memstats":
36                 // Do nothing, this is a big output.
37             default:
38                 log.Printf("\t%s -> %s", kv.Key, kv.Value)
39         }
40     })
41 }
42
43 func main() {
44     flag.Parse()
45     printVars()
46     hello(*times, *name)
47     printVars()
48     hello(*times, *name)
49     printVars()
50 }
```
